Automated Irregular Shaped Stone Stacking via Pose Optimization and Monte Carlo Tree Search

Author Names Omitted for Anonymous Review. Paper-ID [add your ID here]

***Abstract*—Stacking irregular objects into prescribed target shapes without container walls requires mutually supporting arrangements despite uneven contact surfaces. The problem couples combinatorial object selection and ordering with continuous pose optimization, as each placement redistributes loads and changes support geometry and space available for subsequent objects. We propose a framework with two complementary components: structure-aware bilevel pose optimization and parallel Monte Carlo Tree Search (MCTS). The optimizer combines target-shape objectives with contact-force optimization to account for force and torque equilibrium across the structure under friction constraints, using bodywise ADMM updates for GPU parallelization. The receding-horizon MCTS planner evaluates batches of candidate placements in parallel through pose refinement and physical simulation, providing lookahead over object choices and placement poses. We hypothesize that combining structural stability evaluation with lookahead improves target coverage while maintaining stability under a fixed planning budget.**

I. Introduction

Robotic packing methods accommodate irregular objects while considering packing density, stability, and placement feasibility [20, 9]. Constructing freestanding structures from natural stones requires finding mutually supporting arrangements among uneven surfaces while achieving a prescribed target shape without support from container walls.

Several stone-stacking methods use placement criteria developed for particular construction settings, including tower-building objectives [7] and masonry-inspired geometric rules [12, 17]. For example, Johns et al. generate candidates where target wall surfaces meet the existing structure and apply inward-slope heuristics [12]. Extensions to freeform walls [13] broaden wall geometry; more general targets motivate objectives that use target geometry without wall-specific placement rules. Structural stability also requires reasoning beyond the support of the newly placed object: its load can affect the equilibrium of objects beneath it. Liu and Napp [16] explicitly address structural strength during planning, demonstrating the value of evaluating the structure as a whole. However, assessing the current structure alone does not determine how a placement affects future support and space utilization. Menezes et al. [18] address sequential placement through learned action values that account for future rewards, with a prescribed object sequence and fixed orientations. We investigate explicit lookahead over object choices and placement poses, coupled with structural stability evaluation and target-shape objectives.

We consider stacking a finite inventory of n irregular objects to approximate a target shape $T \subset \mathbb{R}^3$. Let $O_i(q_i) \subset \mathbb{R}^3$ denote the volume occupied by object i at pose $q_i \in \mathrm{SE}(3)$.



Fig. 1: Preliminary stacking results. Left: simulated structures for different target shapes. Right: a physical stone structure from an outdoor experiment.

The final configuration objective is

$$\min_{q_1,\cdots,q_n \in \mathrm{SE}(3)} \mathbb{V}\left(T \setminus \bigcup_{i=1}^{n} O_i(q_i)\right) \quad \text{s.t.} \quad c_{\text{phys}}(O_1(q_1),\cdots,O_n(q_n)) = 0, \tag{1}$$

where $\mathbb{V}(\cdot)$ denotes volume and $c_{\text{phys}} = 0$ represents stability and non-interpenetration constraints. Constructing such a configuration requires combinatorial choices of objects and placement order together with continuous pose optimization under coupled physical constraints, making joint planning computationally demanding. We therefore select one object and optimize its placement at each step, checking the resulting structure's physical feasibility. Each placement changes support geometry, free space, and remaining inventory, motivating lookahead beyond immediate placement quality. Our central question is how to combine this lookahead with structural stability evaluation in a common framework for different target shapes.

Our framework couples bilevel pose optimization with batched, receding-horizon Monte Carlo Tree Search (MCTS) [8]. The optimizer combines target geometry with structure-wide force and torque equilibrium under friction constraints, using bodywise ADMM to parallelize contact-force optimization. Different targets share this formulation through their geometric representations. MCTS compares placement sequences using simulation to reject unstable placements and supply settled successor states.

We hypothesize that combining structural stability evaluation with multi-step lookahead improves target coverage while

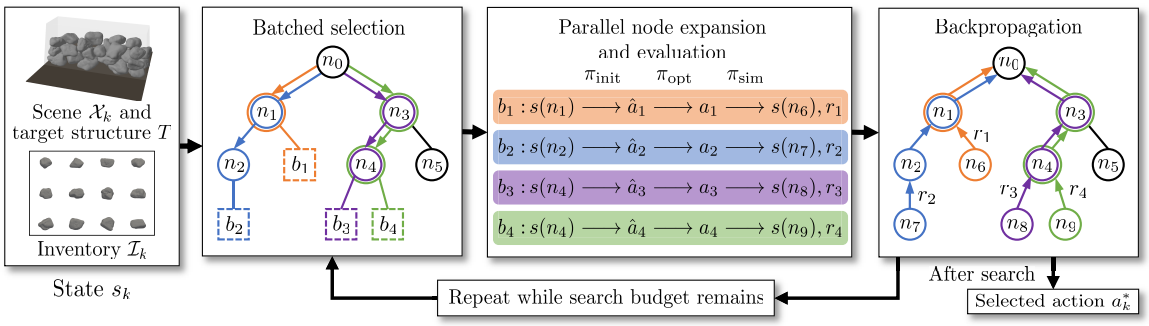


Fig. 2: Batched MCTS sequence planning. Colors identify expansion requests b_j and their tree paths. Requests are evaluated in parallel through π_{init}, π_{opt}, and π_{sim}, and returns are backpropagated along the selected paths. Search repeats within the budget before selecting a physically feasible root action a_k^* from the evaluated tree.

maintaining stability across different target geometries under a fixed planning budget. Comparisons with greedy placement and candidate-only stability evaluation will test the effects of lookahead and structural coupling, respectively. Evaluations across target shapes and search horizons will examine the applicability of the common formulation and its computational tradeoffs. The proposed contributions are:

- A bilevel placement optimizer that combines explicit target geometry with contact-force optimization across the structure, using bodywise ADMM updates for GPU computation.
- A receding-horizon MCTS stacking planner that evaluates batches of candidate placements in parallel through pose refinement and physical simulation, providing multi-step lookahead over object choices and placement poses.

II. Related Works

A. Automated Stone Stacking

Prior stone-stacking methods combine pose search with simulation [7] or use geometric placement rules for wall construction [12, 13, 17]. Liu and Napp [16] incorporate structural strength into sequence planning. Menezes et al. [18] learn placement values with a fixed object sequence and fixed orientations, and their elevation-map strategy informs our initial pose sampler. We combine target-driven pose optimization and structural equilibrium evaluation with explicit lookahead over object choices and placement poses.

B. Contact-Aware Pose Optimization

BODex [5] couples lower-level contact-force optimization with upper-level grasp synthesis, while GraspQP [21] defines a differentiable force-closure energy through a quadratic program. Our optimizer couples candidate placement with structure-wide equilibrium, using bodywise force decomposition for parallel evaluation.

C. Simulation-Based Sequential Manipulation Planning

Parallel Monte Carlo Tree Search with Batched Simulations (PMBS) [8] couples parallel search with GPU rigid-body simulation for object retrieval from clutter. Our stacking planner adopts batched evaluation, integrating pose refinement and structural simulation at each expansion.

III. Preliminaries and Overview

A. Shape Representation

We represent the target T and each object O_i as unions of smooth convex bodies modeled by differentiable support functions (DSFs) [14]. For a compact convex body $C \subset \mathbb{R}^3$, its support function is $h_C(x) = \max_{p \in C} x^\top p$ for $x \neq 0$. At pose $q = (R, t) \in \mathrm{SE}(3)$, we write

$$C(q) = \{Rp + t \mid p \in C\}, \qquad O_i(q_i) = \bigcup_j C_{ij}(q_i),$$

where C_{ij} are the convex components in object i's local frame. Similarly, $T = \bigcup_j T_j$, where T_j are fixed convex components in the world frame. DSFs support differentiable signed-distance and contact-geometry computations between posed convex bodies [2], enabling gradient-based pose optimization.

B. Markov Decision Process Formulation

To address the construction objective in Eq. (1), we model stacking as a finite-horizon sequential decision process in which one object is placed at each step. Let $\mathcal{P}_k \subseteq \{1, \dots, n\}$ be the indices of the objects already placed, and let $i_0 \notin \{1, \dots, n\}$ index the fixed ground. The current scene is $\mathcal{X}_k = \{(i, q_i) \mid i \in \mathcal{P}_k\}$, and the remaining inventory is $\mathcal{I}_k = \{1, \dots, n\} \setminus \mathcal{P}_k$. The state, action, and deterministic transition function are

$$s_k = (\mathcal{X}_k, \mathcal{I}_k, T), \quad a_k = (i_{k+1}, q_{k+1}), \quad s_{k+1} = F(s_k, a_k),$$

where $i_{k+1} \in \mathcal{I}_k$. The transition inserts the proposed pose into the scene, simulates the resulting structure, and returns the settled scene with $\mathcal{I}_{k+1} = \mathcal{I}_k \setminus \{i_{k+1}\}$. The reward $r(s_k, a_k)$ scores target-contained volume and physical quality; feasibility checks separately reject invalid transitions (Sec. V). The sequence planner therefore operates on the MDP $\mathcal{M} = (\mathcal{S}, \mathcal{A}, F, r, \gamma)$.

C. Monte Carlo Tree Search

Monte Carlo Tree Search (MCTS) incrementally constructs a search tree $\mathcal{T} = (\mathcal{V}, \mathcal{E})$ over the MDP. Each node $n \in \mathcal{V}$ represents a state $s(n)$, and each edge $(n, a) \in \mathcal{E}$ represents an action and its transition to a child node. Let $N(n)$ denote the number of completed evaluations and $Q(n)$ the empirical mean of discounted returns backpropagated to that node.

Each iteration selects a path by balancing value and exploration, expands and evaluates an action, and backpropagates its discounted return to update Q and N.

In our planner (Fig. 2), expansion uses three modules: π_{init} proposes an object and initial pose $\hat{a}_k = (i_{k+1}, \hat{q}_{k+1})$, π_{opt} refines the pose, and π_{sim} evaluates the resulting structure to determine feasibility, the settled successor, and reward. The resulting flow is

$$s_k \xrightarrow{\pi_{\text{init}}} \hat{a}_k \xrightarrow{\pi_{\text{opt}}} a_k \xrightarrow{\pi_{\text{sim}}} s_{k+1},\ r(s_k, a_k).$$

Sec. V describes how these modules are integrated into batched search.

IV. Pose Optimization

The module π_{opt} refines $\hat{q}_{k+1}$ from π_{init} toward a stable, physically feasible placement q_{k+1}.

A. Single Object Pose Optimization Problem

Given s_k, we initialize q at $\hat{q}_{k+1}$ and solve the bilevel problem

$$\begin{aligned} \min_{q} \quad & \left[f_{\text{geo}}(q) + \min_{\lambda} f_{\text{phys}}(q, \lambda) \right] \\ \text{s.t.} \quad & g_{\text{cone},j}(q, \lambda_j) \le 0, \quad j = 1, \cdots, n_c \end{aligned} \tag{2}$$

where $q \in \text{SE}(3)$ is the candidate object's pose and $\lambda = (\lambda_1, \cdots, \lambda_{n_c}) \in \mathbb{R}^{3n_c}$ collects all contact forces in the resulting structure, including inter-object and ground contacts. This formulation follows prior contact-aware optimization methods [15, 5].

The geometric objective f_{geo} is designed as follows:

$$\begin{aligned} f_{\text{geo}}(q) = w_{\text{h}} f_{\text{h}}(q) + w_{\text{in}} f_{\text{in}}(q \mid T, O_{i_{k+1}}) \\ + w_{\text{pen}} f_{\text{pen}}(q \mid \mathcal{X}_k, O_{i_{k+1}}) \end{aligned}$$

with positive weights. For $q = (R, t)$, $f_{\text{h}}(q) = t_z$ is the world-frame height of the candidate's pose origin. Writing $O = O_{i_{k+1}}$, we penalize proximity to existing objects and the fixed ground as

$$f_{\text{pen}}(q) = \sum_{j \in \mathcal{P}_k \cup \{i_0\}} \left[\epsilon_{\text{margin}} - \text{d}(O(q), O_j(q_j)) \right]_+^2,$$

where d is signed distance (positive for separation), $[x]_+ = \max(x, 0)$, and $\epsilon_{\text{margin}} \ge 0$ is the clearance margin. With $h_{O(q)}(u) = \max_j h_{C_{i_{k+1},j}(q)}(u)$, we define

$$\begin{aligned} m_{\text{in}}(q) &= \max_j \min_{u \in \mathbb{S}^2} \left[h_{T_j}(u) - h_{O(q)}(u) \right], \\ f_{\text{in}}(q) &= \left[-m_{\text{in}}(q) \right]_+^2. \end{aligned}$$

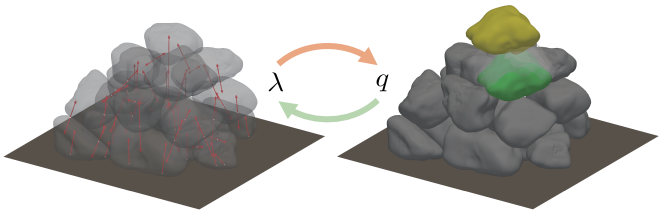


Fig. 3: Bilevel pose optimization. Left: optimizing contact forces λ under friction constraints minimizes structural equilibrium residuals at fixed q. Right: refining $q \in \text{SE}(3)$ minimizes the geometric and optimized physical objectives.

This term penalizes violations of containment of the entire object within a single convex target component, a conservative condition for containment in T. Noting that the term is differentiable almost everywhere in q, we compute pose derivatives at regular inner solutions by implicit differentiation. Nonsmooth transitions have not caused practical difficulties in our implementation.

Let $\mathcal{P}_{k+1} := \mathcal{P}_k \cup \{i_{k+1}\}$. For distinct $l, m \in \mathcal{P}_{k+1} \cup \{i_0\}$, let $I_{l,m} \subset \{1, \ldots, n_c\}$ index contacts between bodies l and m. Each object's contact set $I_l := \bigcup_{m \in (\mathcal{P}_{k+1} \cup \{i_0\}) \setminus \{l\}} I_{l,m}$ includes ground contacts, but the physical objective excludes the ground's equilibrium residual:

$$\begin{aligned} f_{\text{phys}}(q, \lambda) = w_{\text{eq}} \sum_{l \in \mathcal{P}_{k+1}} \| J_{c,I_l}^{\top}(q) \lambda_{I_l} + \lambda_{\text{ext},l} \|_W^2 \\ + w_{\text{cp}} \sum_{j \in I_{i_{k+1}}} \text{d}_j(q)^2 \|\lambda_j\|^2 + w_{\text{reg}} \sum_{j=1}^{n_c} \|\lambda_j\|^2 \end{aligned} \tag{3}$$

The three terms respectively penalize violations of static equilibrium, contact complementarity between the candidate and the current scene, and large contact forces through regularization. Figure 4 illustrates why equilibrium is evaluated across the structure rather than only for the candidate. Here, $\text{d}_j(q)$ is the signed distance between the two bodies associated with contact candidate j. J_{c,I_l} is the contact Jacobian for object O_l, $\lambda_{\text{ext},l}$ is its external wrench, λ_{I_l} is the vector of contact forces at the contact points in I_l, and $w_{\text{eq}}, w_{\text{cp}}, w_{\text{reg}} > 0$ are objective weights.

Finally, each $g_{\text{cone},j}(\cdot)$ is the smoothed friction cone constraints for the j-th contact point defined by

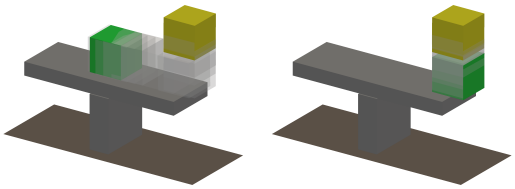
$$g_{\text{cone},j}(q, \lambda_j) = \sqrt{\|J_{c,j}^t \lambda_j\|^2 + \epsilon} - \mu_j J_{c,j}^n \lambda_j \tag{4}$$

where $\epsilon > 0$ is a smoothing parameter, μ_j is the friction coefficient, and $J_{c,j}^t$ and $J_{c,j}^n$ are the tangential and normal components of the Jacobian of the contact constraint.

B. Upper-Level Problem Optimization on SE(3) Manifold

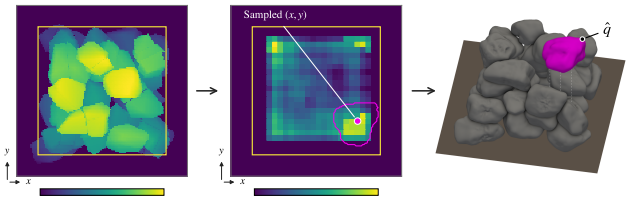
Let $\lambda^*(q)$ be the solution of the lower-level problem and define $f_{\text{phys}}(q) := f_{\text{phys}}(q, \lambda^*(q))$. The envelope theorem [4] gives

$$\frac{df_{\text{phys}}(q)}{dq} = \frac{\partial f_{\text{phys}}(q, \lambda^*)}{\partial q} + \sum_{j=1}^{n_c} \nu_j^* \frac{\partial g_{\text{cone},j}(q, \lambda_j^*)}{\partial q}$$



(a) Structure-wide (b) Candidate-only

Fig. 4: Seesaw example: initial (yellow), optimized (green), and intermediate (gray) poses. Structure-wide equilibrium (a) moves the candidate toward the support; candidate-only equilibrium (b) leaves it near the unsupported end.



(a) Scene height (b) Proposal score (c) Sampled initial pose

Fig. 5: Initial pose sampling for a fixed object and orientation: (a) scene height, (b) proposal score and sampled planar location, and (c) initial pose $\hat{q}$ above the scene. Magenta links the sampled location and candidate object.

where ν_j^* is the optimal dual variable for the j-th friction cone constraint. Using the above gradient, we optimize the placement pose q on SE(3) manifold based on the Riemannian trust-region method [1] combined with BFGS Hessian approximation [10]. A homotopy schedule adjusts the geometric and physical weights between stages, holding them fixed within each stage and initializing from the preceding pose. It progressively reduces w_{reg} to limit regularization-induced equilibrium residuals. Candidate evaluation uses the final pose and forces.

C. Lower-Level Problem Optimization with ADMM

For fixed q, the lower-level problem is convex in λ. To apply bodywise ADMM [3], we introduce two local copies $\lambda_{j,1}$ and $\lambda_{j,2}$ per contact, one for each contacting body, including the ground. Collecting these copies into $\tilde{\lambda}$ makes each body's equilibrium residual depend only on its local variables. Consensus and friction constraints give the reformulation

$$\begin{aligned} \min_{\tilde{\lambda},z} \quad & \tilde{f}_{\text{phys}}(q,\tilde{\lambda}) + \sum_{\substack{j=1,\dots,n_c \\ s=1,2}} \delta_{K_j}(z_{j,s}) \\ \text{s.t.} \quad & \lambda_{j,1} = \lambda_{j,2} = z_{j,1} = z_{j,2}, \qquad j = 1,\dots,n_c. \end{aligned}$$

where $\tilde{f}_{\text{phys}}$ is the physical objective in local variables, $z_{j,s}$ are auxiliary variables, and δ_{K_j} is the indicator of $K_j := \{\lambda \in \mathbb{R}^3 \mid g_{\text{cone},j}(q,\lambda) \le 0\}$. ADMM solves the body subproblems independently. An equivalent four-dimensional second-order cone (SOC) lifting enables closed-form cone projections for the friction updates, while consensus updates reduce to averaging.

The pose optimization module is implemented using CUDA to leverage GPU parallelization for efficient computation of the gradients and the ADMM updates.

V. Sequence Planning

A. Initial Pose Sampling

The sampler π_{init} uses volume-biased object sampling with a finite orientation set per object. It favors lower placements and seeds empty scenes near feasible target corners. For object i, orientation R, and feasible planar location $\mathbf{x}$, the score is

$$\sigma(\mathbf{x} \mid i,R) = w_{\text{s}}\sigma_{\text{s}} + w_{\text{f}}\sigma_{\text{f}} + w_{\text{b}}\sigma_{\text{b}} - w_{\text{h}}\sigma_{\text{h}} - w_{\text{d}}\sigma_{\text{d}},$$

with positive weights. The first four terms measure the fraction of bottom samples near predicted contact, unfilled layer area beneath the footprint, proximity to the feasible placement boundary, and normalized mean contact height, respectively. A bounded penalty σ_{d} on the planar footprint gap to the preceding object favors nearby placements, helping MCTS assess interactions between successive objects without excluding distant proposals. It is zero for the first placement. The floor and boundary preferences are inspired by Menezes et al. [18]. Figure 5 illustrates the proposal score and sampled initial pose.

We combine elite selection with sampling without replacement weighted by $\exp(\sigma/\tau)$, $\tau > 0$. Proposals start above the scene with clearance and undergo π_{opt} refinement, including XY/yaw mutations at the root. Heuristic filters limit planar displacement from each initial pose after optimization and settling, and screen out candidates whose bottoms lie too far above the current construction level.

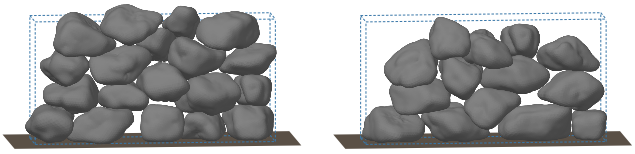
B. Simulation Module

The module π_{sim} evaluates every proposed transition using the same simulation procedure at all search depths. It first settles the new object for a fixed duration while holding the existing objects static. From the resulting poses, it resets velocities and simulates the entire structure's response to the added load. Structural simulation requires a minimum duration and sustained low mass-normalized kinetic energy. Excessive accumulated motion, invalid contacts, or failure to settle within the time limit rejects the placement.

For an accepted nominal simulation, the settled scene defines the successor $s' = F(s,a)$. Each disturbance trial starts from this scene with the new object's initial world-frame linear velocity

$$\xi_j = \left(v_h \cos\theta_j,\ v_h \sin\theta_j,\ -v_d\right)^\top,$$

where $\theta_j = 2\pi j/n_d$ for $j = 0,\dots,n_d-1$, and $v_h, v_d > 0$ are prescribed horizontal and downward speeds. Initial angular velocity is zero, and every candidate uses the same n_d disturbances. To reduce computation, only the candidate and stones within two links in the nominal scene's contact graph of movable stones remain dynamic throughout each trial. All other bodies retain fixed collision geometry. These trials



(a) MCTS (b) Greedy

Fig. 6: Wall construction from the same 49-stone inventory (seed 6): MCTS places 19 stones (44.0% target filling); greedy stops at 14 (36.1%), finding no further feasible placement. Blue dashed outlines show the target.

estimate local robustness without changing s'. Simulation uses CUDA subADMM [11].

C. Reward Design

For the newly placed object $i = i_{k+1}$, let q_i' be its pose in the settled successor s' returned by π_{sim} (Sec. V-B). We define the target-volume and object-volume ratios

$$\rho_T = \frac{\mathbb{V}\big(O_i(q_i') \cap T\big)}{\mathbb{V}(T)}, \quad \rho_O = \frac{\mathbb{V}\big(O_i(q_i') \cap T\big)}{\mathbb{V}(O_i)}.$$

The ρ_T term encourages target filling toward the objective in Eq. (1), while ρ_O discourages placements extending excessively outside the target. Intersection volumes are estimated on a voxel grid. With fixed normalization scales absorbed into positive weights, the standard excavator reward is

$$\begin{aligned} r(s,a) = w_T\rho_T + w_O[\rho_O - \tfrac{1}{2}]_+ - w_E e(a) \\ + w_H\big(\phi_H(s) - \phi_H(s')\big) - w_D D, \end{aligned}$$

where $[x]_+ = \max(x, 0)$, so the containment bonus is active above 50%. Coverage lost through motion of earlier objects is not subtracted. The height term rewards the decrease in surface-height variance $\phi_H(s) = \langle (H_s - \langle H_s \rangle_\Omega)^2 \rangle_\Omega$ between parent and settled successor. Here H_s is height over the target footprint Ω, zero at empty cells, and $\langle\cdot\rangle_\Omega$ denotes the area-weighted mean.

The equilibrium score uses residual forces $\mathbf{f}_l$ and torques $\boldsymbol{\tau}_l$ returned by pose optimization:

$$e(a) = \min\left\{1, \max_l \left(\frac{\|\mathbf{f}_l\|}{m_l g} + \frac{\|\boldsymbol{\tau}_l\|}{m_l g \mathbb{V}(O_l)^{1/3}}\right)\right\}.$$

Here l indexes objects in the resulting structure and g is gravitational acceleration. Let $q_l(t; \xi_j)$ denote object l's pose under the predefined velocity disturbance ξ_j (Sec. V-B), starting from its nominal successor pose q_l'. Disturbance sensitivity is

$$D = \min\left\{D_{\max}, \max_{j,l,t} d_{\text{SE(3)}}\big(q_l(t; \xi_j), q_l'\big)\right\},$$

where $d_{\text{SE(3)}}$ sums Euclidean translation distance and angular distance weighted by $\lambda_\theta > 0$. The maximum spans predefined disturbances, objects l in this local dynamic neighborhood, and simulation time; $D_{\max}$ caps the penalty. Physical rejection criteria apply independently of the reward.

D. Monte Carlo Tree Search

We integrate these modules in batched, receding-horizon MCTS (Fig. 2), following the parallel evaluation approach of [8]. Figure 6 shows a recorded comparison with greedy placement.

Batched Selection. For an optimized action $a = (i, q)$, let ρ_O^{pre} denote the object-volume ratio evaluated at q, before simulation. Using the reward weights, the logit $\ell(a) = w_O\rho_O^{\text{pre}} - w_E e(a)$ defines $P(n_p, a)$ through a softmax over selectable children. Writing $N_n = N(n)$ and similarly for Q_n and $\hat{N}_n$, an evaluated child n of n_p is selected using PUCT [19]:

$$\text{PUCT}(n) = \frac{N_n Q_n}{N_n + \hat{N}_n} + cP(n_p, a)\frac{\sqrt{N_{n_p} + \hat{N}_{n_p}}}{1 + N_n + \hat{N}_n}, \quad (5)$$

where $c > 0$ controls exploration, $\hat{N}$ counts temporary reservations, and the first term is zero when its denominator is zero. At non-root nodes, progressive widening [6] increases the number of candidate actions with completed visits. Temporary reservations prevent duplicate evaluations within a batch without increasing the widening allowance.

Parallel Expansion and Evaluation. After batch selection, fresh proposals pass through $\pi_{\text{init}} \to \pi_{\text{opt}} \to \pi_{\text{sim}}$. Independent candidates are processed in parallel; selection never descends through an action without an evaluated successor. Failed candidates are excluded from further selection.

Leaf Value Estimation and Backpropagation. We replace rollouts with a discounted per-step reward heuristic, capped by depth and remaining inventory. Terminal, failed, and depth-limit leaves receive no bootstrap; exhausted shallower nonterminal leaves are not repeatedly backed up. Failures incur a discounted fixed per-step penalty over the remaining horizon, capped by inventory. Backpropagation updates Q and N and clears temporary reservations.

Action Selection and Tree Reuse. After search, we select an action using only the evaluated tree. Let d be the required number of descendant edges below a root candidate, capped by the remaining search horizon. For a root action a, let $G_d(a)$ be the maximum discounted return over existing feasible paths reaching this depth, including the reward of a. Among root actions with such a path, we maximize

$$J(a) = (1 - \beta)Q(n_a) + \beta G_d(a), \qquad \beta \in [0, 1].$$

The path-return term counterbalances failures accumulated during exploration, which can depress the mean even when a promising feasible path exists. If no root action meets the depth requirement, we select the feasible root action with the highest mean return, or return no placement if none is feasible.

After execution, we advance the root to the observed successor. Evaluated descendants are retained when this state and the planning context match those used to construct the subtree; otherwise, the affected subtree and its statistics are discarded.

Acknowledgments

References

[1] P. A. Absil, C. G. Baker, and K. A. Gallivan. Trust-region methods on riemannian manifolds. *Foundations of Computational Mathematics*, 7:303–330, 2007.

[2] S. An, S. Lee, J. Lee, S. Park, and D. Lee. Collision detection between smooth convex bodies via riemannian optimization framework. In *2024 IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*, pages 12464–12471. IEEE, 2024.

[3] S. Boyd, N. Parikh, E. Chu, B. Peleato, and J. Eckstein. Distributed optimization and statistical learning via the alternating direction method of multipliers. *Foundations and Trends in Machine Learning*, 3(1):1–122, 2011. doi: 10.1561/2200000016.

[4] M. Carter. *Foundations of mathematical economics*. MIT press, 2001.

[5] J. Chen, Y. Ke, and H. Wang. BODex: Scalable and efficient robotic dexterous grasp synthesis using bilevel optimization. In *2025 IEEE International Conference on Robotics and Automation (ICRA)*, pages 1–8. IEEE, 2025. doi: 10.1109/ICRA55743.2025.11127930.

[6] A. Couëtoux, J.-B. Hoock, N. Sokolovska, O. Teytaud, and N. Bonnard. Continuous upper confidence trees. In *Learning and Intelligent Optimization*, volume 6683 of *Lecture Notes in Computer Science*, pages 433–445. Springer, 2011. doi: 10.1007/978-3-642-25566-3_32.

[7] F. Furrer, M. Wermelinger, H. Yoshida, F. Gramazio, M. Kohler, R. Siegwart, and M. Hutter. Autonomous robotic stone stacking with online next best object target pose planning. In *2017 IEEE International Conference on Robotics and Automation (ICRA)*, pages 2350–2356, 2017. doi: 10.1109/ICRA.2017.7989272.

[8] B. Huang, A. Boularias, and J. Yu. Parallel monte carlo tree search with batched rigid-body simulations for speeding up long-horizon episodic robot planning. In *2022 IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*. IEEE, 2022.

[9] S. Huang, Z. Wang, J. Zhou, and J. Lu. Planning irregular object packing via hierarchical reinforcement learning. *IEEE Robotics and Automation Letters*, 8(1): 81–88, 2023. doi: 10.1109/LRA.2022.3222996.

[10] W. Huang, P.-A. Absil, and K. A. Gallivan. A riemannian BFGS method without differentiated retraction for nonconvex optimization problems. *SIAM Journal on Optimization*, 28(1):470–495, 2018. doi: 10.1137/17M1127582.

[11] H. Ji, H. Kim, J. Lee, S. Lee, S. An, J. Heo, Y. Lee, Y. Lee, and D. Lee. Gpu-accelerated subsystem-based admm for large-scale interactive simulation. In *2025 IEEE International Conference on Robotics and Automation (ICRA)*, pages 13868–13874. IEEE, 2025.

[12] R. L. Johns, M. Wermelinger, R. Mascaro, D. Jud, F. Gramazio, M. Kohler, M. Chli, and M. Hutter. Autonomous dry stone. *Construction Robotics*, 4(3–4):127–140, 2020. doi: 10.1007/s41693-020-00037-6.

[13] R. L. Johns, M. Wermelinger, R. Mascaro, D. Jud, I. Hurkxkens, L. Vasey, M. Chli, F. Gramazio, M. Kohler, and M. Hutter. A framework for robotic excavation and dry stone construction using on-site materials. *Science Robotics*, 8(84):eabp9758, 2023. doi: 10.1126/scirobotics.abp9758.

[14] J. Lee, M. Lee, and D. Lee. Uncertain pose estimation during contact tasks using differentiable contact features. *RSS 2023-Robotics: Science and Systems*, 2023.

[15] J. Lee, S. Park, M. Lee, and D. Lee. Efficient gradient-based inference for manipulation planning in contact factor graphs. In *2025 IEEE International Conference on Robotics and Automation (ICRA)*, pages 15401–15407. IEEE, 2025. doi: 10.1109/ICRA55743.2025.11127769.

[16] Y. Liu and N. Napp. Stability-based sequence planning for robotic dry-stacking of natural stones. *IEEE Robotics and Automation Letters*, 8(9):5894–5901, 2023. doi: 10.1109/LRA.2023.3301275.

[17] Y. Liu, J. Choi, and N. Napp. Planning for robotic dry stacking with irregular stones. In G. Ishigami and K. Yoshida, editors, *Field and Service Robotics*, volume 16 of *Springer Proceedings in Advanced Robotics*, pages 321–335. Springer, 2021. doi: 10.1007/978-981-15-9460-1_23.

[18] A. Menezes, P. Vicente, A. Bernardino, and R. Ventura. From rocks to walls: A model-free reinforcement learning approach to dry stacking with irregular rocks. In *2021 IEEE/CVF Conference on Computer Vision and Pattern Recognition Workshops (CVPRW)*, pages 2057–2065, 2021. doi: 10.1109/CVPRW53098.2021.00234.

[19] D. Silver, J. Schrittwieser, K. Simonyan, I. Antonoglou, A. Huang, A. Guez, T. Hubert, L. Baker, M. Lai, A. Bolton, Y. Chen, T. Lillicrap, F. Hui, L. Sifre, G. van den Driessche, T. Graepel, and D. Hassabis. Mastering the game of go without human knowledge. *Nature*, 550:354–359, 2017.

[20] F. Wang and K. Hauser. Dense robotic packing of irregular and novel 3D objects. *IEEE Transactions on Robotics*, 38(2):1160–1173, 2022. doi: 10.1109/TRO.2021.3097261.

[21] R. Zurbrügg, A. Cramariuc, and M. Hutter. GraspQP: Differentiable optimization of force closure for diverse and robust dexterous grasping. In *Conference on Robot Learning (CoRL)*, pages 2583–2602. PMLR, 2025.